


Popular Applications



Applications & Architectures

A sample set of cloud applications

- Resource usage model
- Challenges and opportunities while moving towards serverless cloud
- System architecture, programming model, design and deployment framework

<i>Application</i>	<i>Description</i>
<i>Feature-rich Scalable Web Application</i>	<i>Microservices Architecture</i>
<i>Real-Time Video</i>	<i>On-the-fly video encoding</i>
<i>Unstructured Data Processing</i>	<i>Big data processing on large clusters</i>
<i>Linear Algebra</i>	<i>Large scale linear algebra</i>
<i>ML Pipelines</i>	<i>ML Training at scale</i>
<i>Databases</i>	<i>OLTP</i>

Application Architectures

- Defines the relation of different components of the system
(application, middleware, data stores) that work together for an application

[Traditional] • N-Tier

[Popular] • Microservices

- Big Compute
- Big Data
- Event-Driven
- Web-Queue-Worker

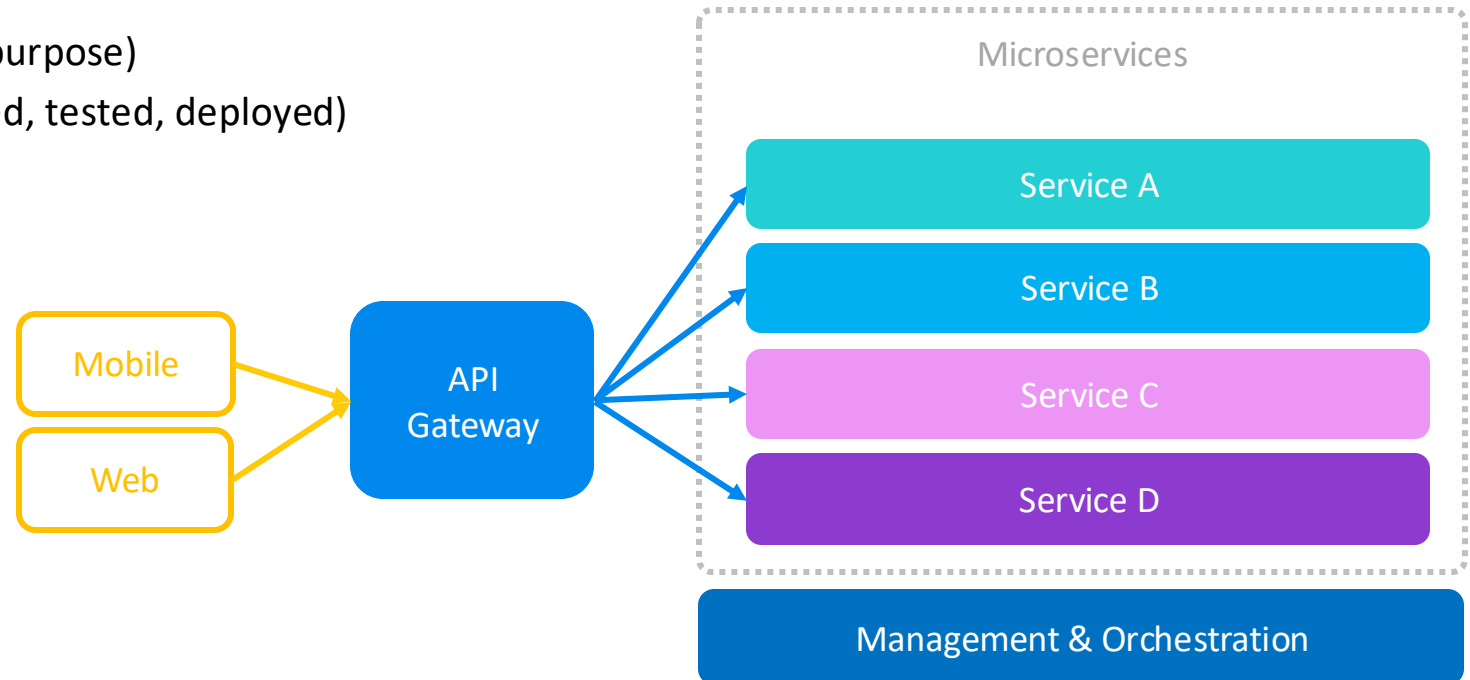
Relation?

Event-Bus Broker Blackboard
Primary-Secondary Interpreter
MVVM Monolithic Pipe-Filter
Client-Server Peer to Peer
Model-View-ViewModel
Layered Model-View-Controller

Different architectures could be implemented using different frameworks and technologies

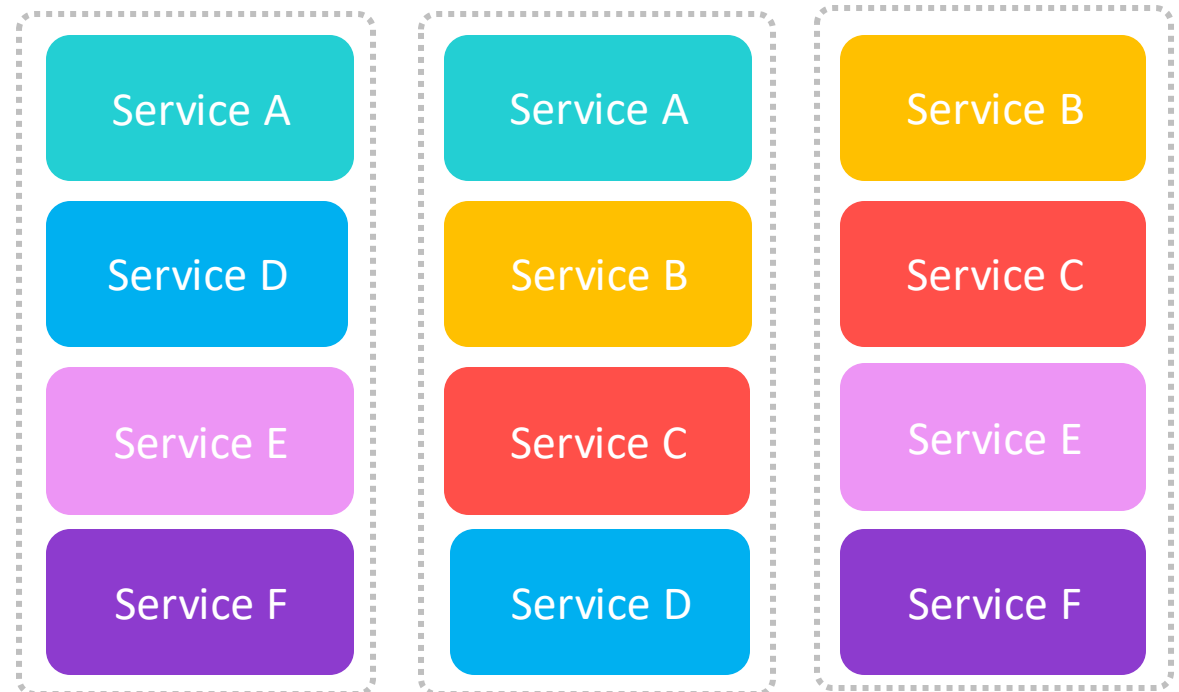
Microservices

- Context: Server-side enterprise application
- Characteristics
 - Small (Cohesion around business purpose)
 - Autonomous (Separately developed, tested, deployed)
 - Single responsibility
 - Independent of other services



Microservices

- Each service is a separate codebase
- Not shared
 - Technology stack
 - Libraries
 - Frameworks
- Hide details from other services
- Services can be deployed independently
 - No application rebuilt

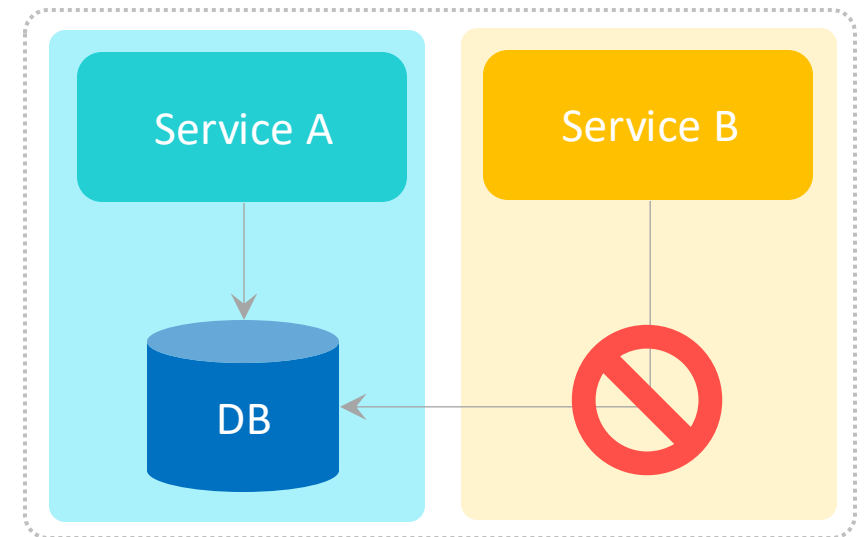


Coupling

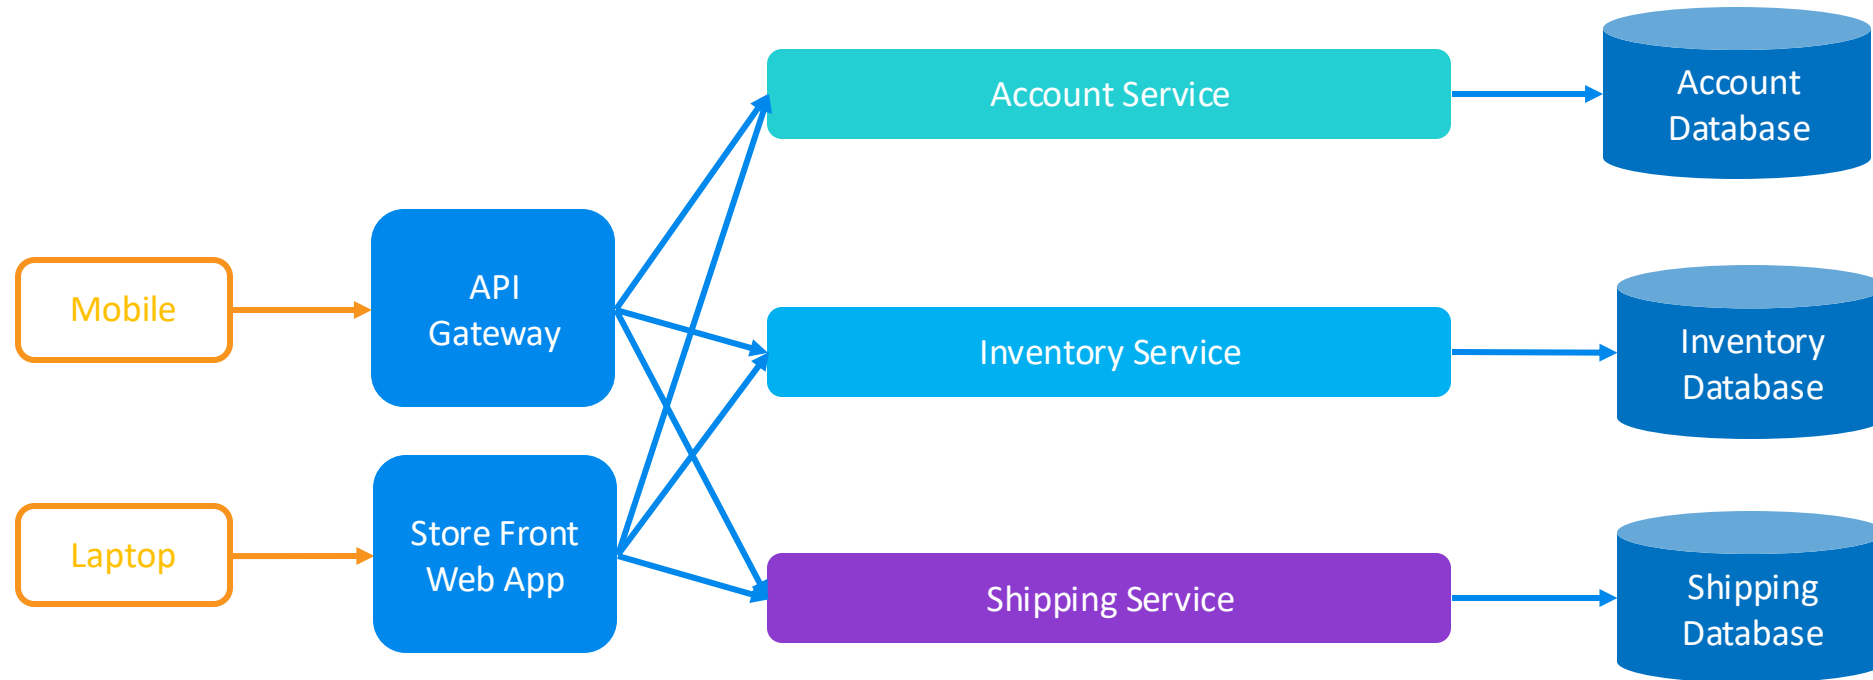
- Domain Coupling
 - One Microservice the user of the functionality of the other Microservice
- Pass through Coupling
 - One microservice passes data to another microservice because it is needed by some other further downstream microservice
- Common Coupling
 - Two or more microservices make use of a common set of data
- Content Coupling
 - An upstream service reaches into the internals of a downstream service and changes its internal state
- Temporal Coupling
 - One microservice needs to call another microservice in a synchronous way

Microservices

- Services are responsible for persisting their own data or external state
 - Data is private to the service that owns it
 - Services do not share data
 - No data persistence layer
 - Leads to polyglot persistence
- Challenges
 - Redundancy
 - Same data in multiple places
 - Tracking Updates (Eventual Consistency)
 - Use events (aggregation & batching)
 - What if strong consistency needed? One service source



Microservices: Example



Examples

- Well-Known Live Services

- Amazon
- Netflix
- eBay
- LinkedIn
- Uber

Code examples

Azure Drone Delivery Example

<https://github.com/mspnp/microservices-reference-implementation>

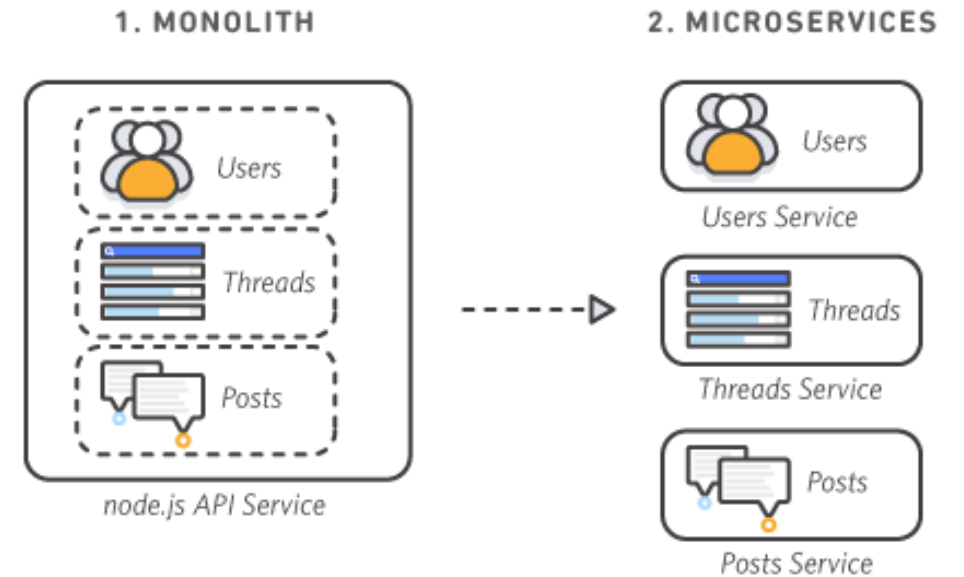
Oracle microservices deployment K8 example

<https://docs.oracle.com/en/solutions/deploy-microservices/index.html>

<https://github.com/oracle-devrel/terraform-oci-arch-microservice-oke>

Additional Examples

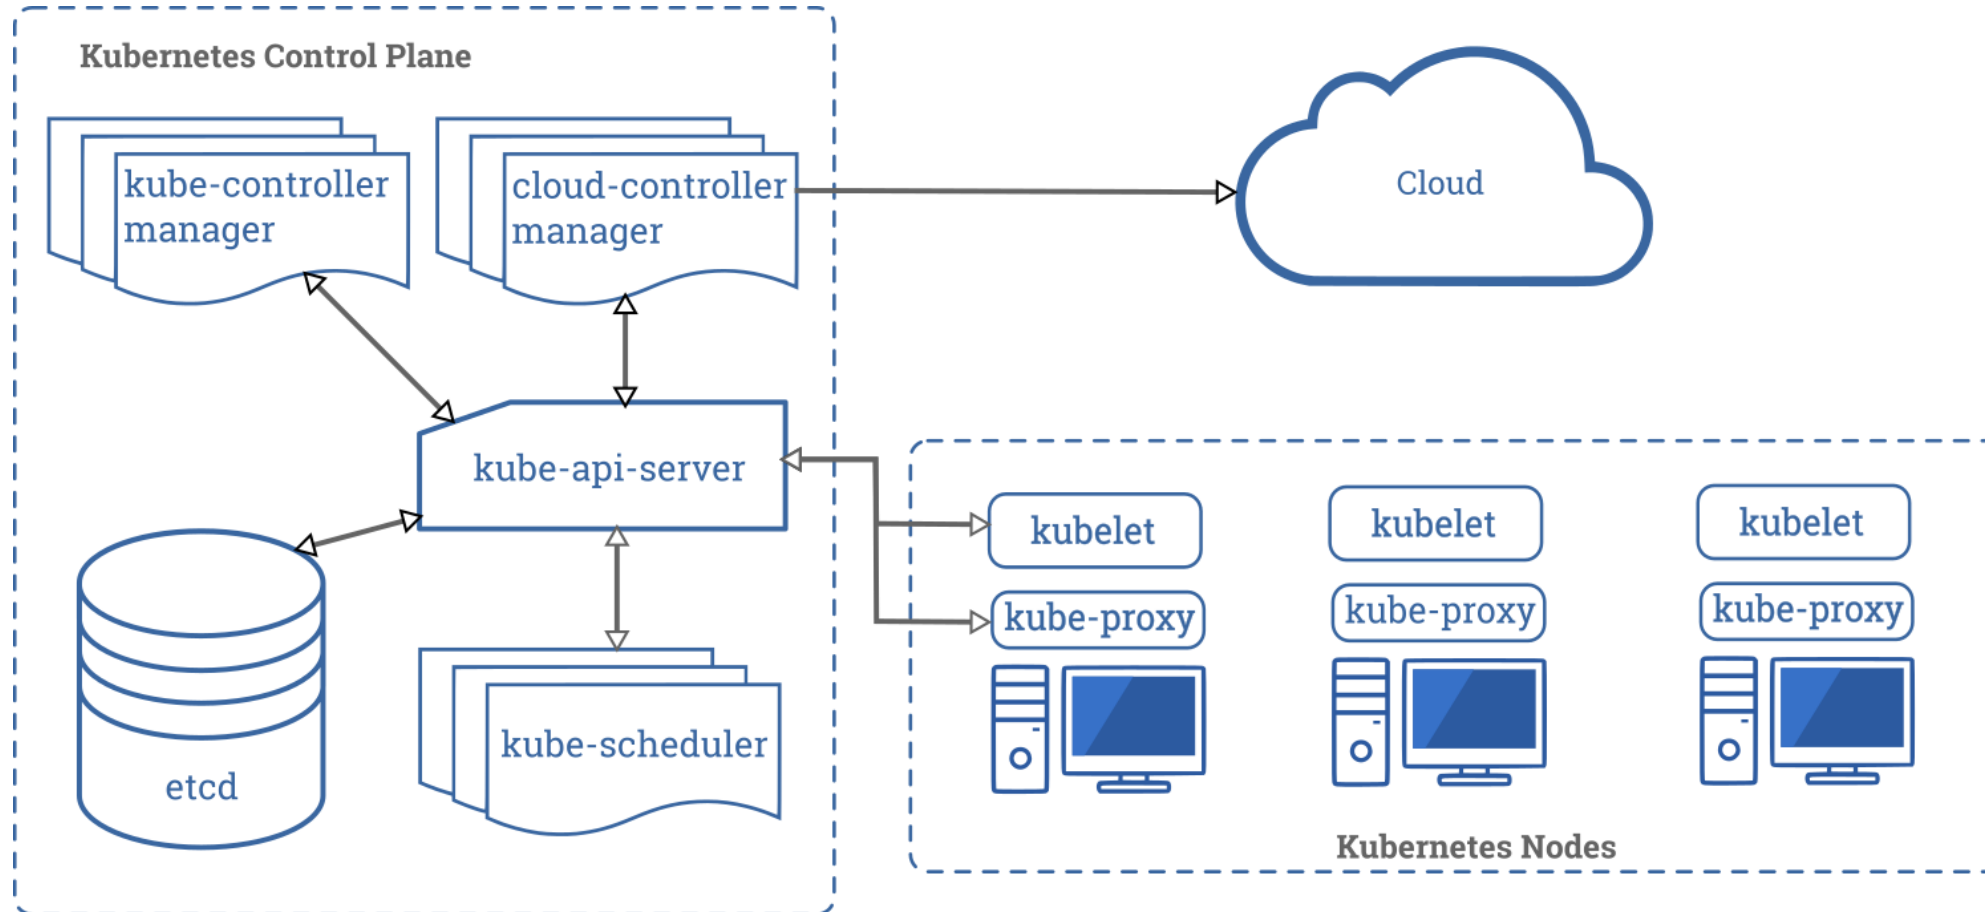
<https://eventuate.io/exampleapps.html>



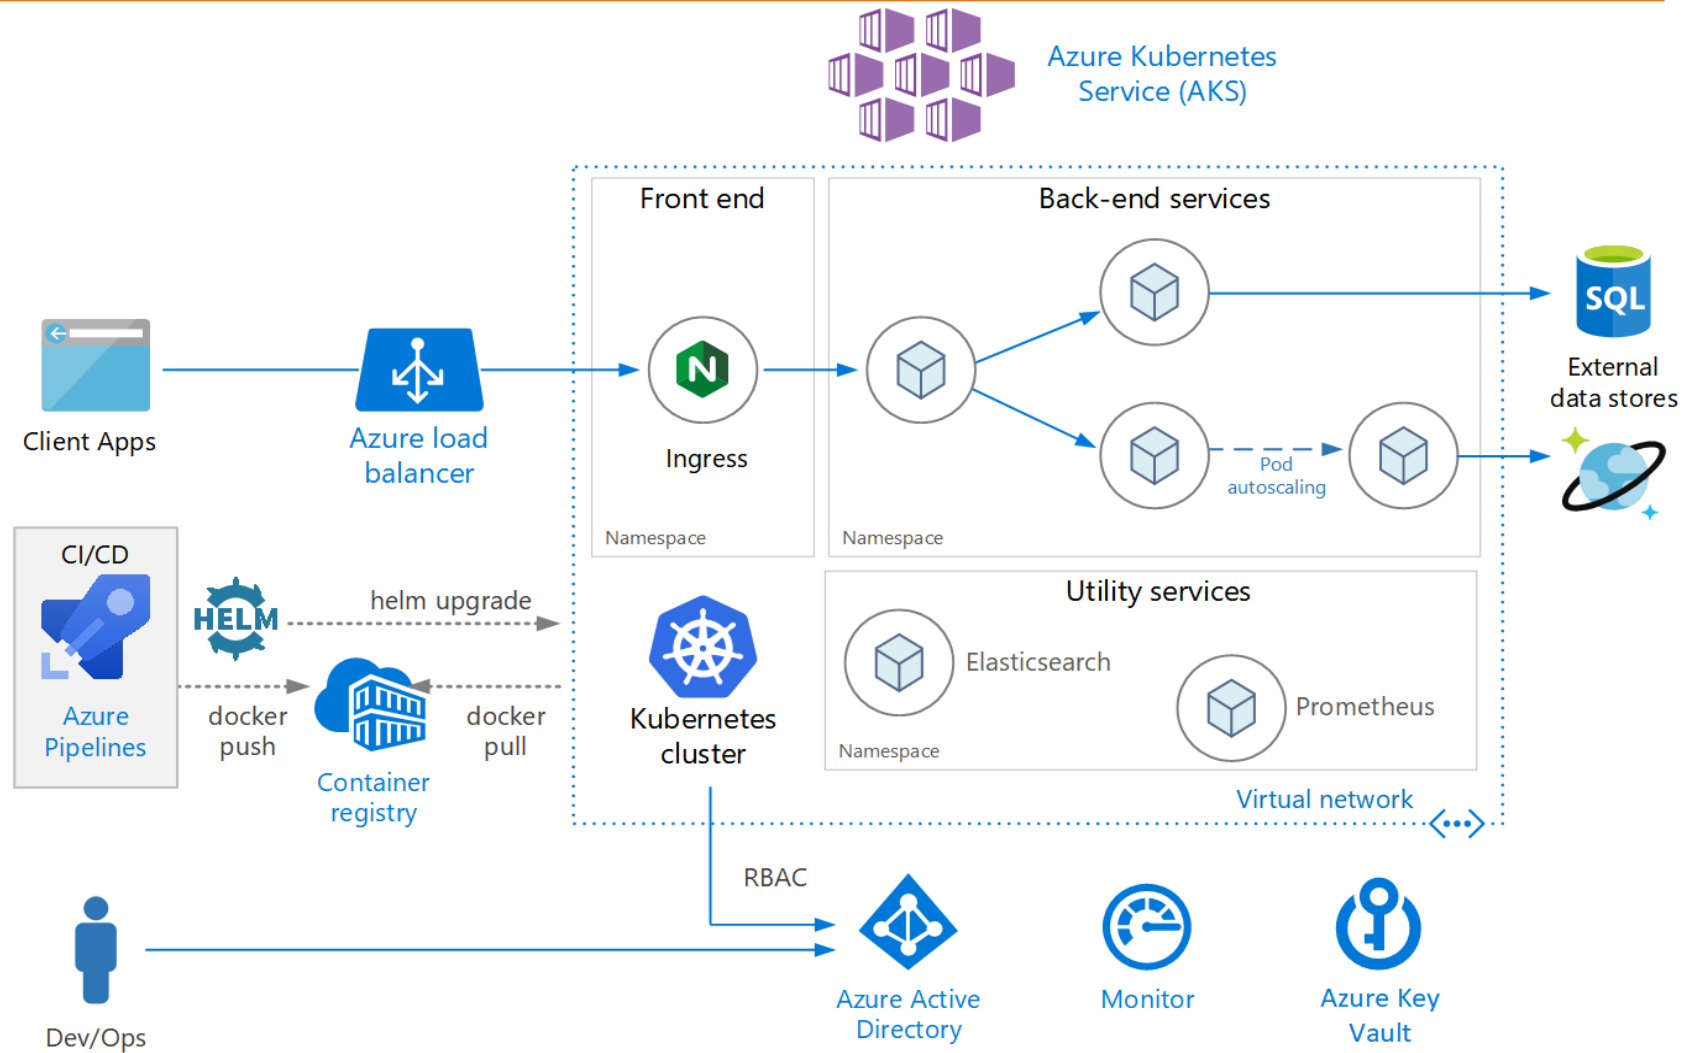
Kubernetes (Greek word meaning helmsman or pilot)

- Open-source platform for managing containerized workloads
- Kubernetes Components
 - Master (Control Plane)
 - API Server
 - etcd: Consistent and highly-available key value store
 - Kube scheduler: Pod to node assignment
 - Kube-controller manager: Control Plane component that runs controller processes
 - Cloud-controller manager: Runs controllers that interact with the underlying cloud providers
 - Nodes (Minions): A worker machine (May be a VM or physical machine, depending on the cluster)
 - Container (Most often Docker)
 - Kubelet: An agent that runs on each node in the cluster and makes sure that containers are running in a pod
 - Kube-proxy: A network proxy that runs on each node in your cluster
 - Workloads
 - Pod: Basic execution unit of a Kubernetes application
 - Controllers: Controller can create and manage multiple Pods


Deployment: Kubernetes



Deployment: K8



Serverless



Knative

- Kubernetes Extension with set of middleware components
- Kubernetes-native APIs for deploying serverless-style functions

- Components
 - Eventing: Management and delivery of events
 - Serving: Request-driven compute that can scale to zero



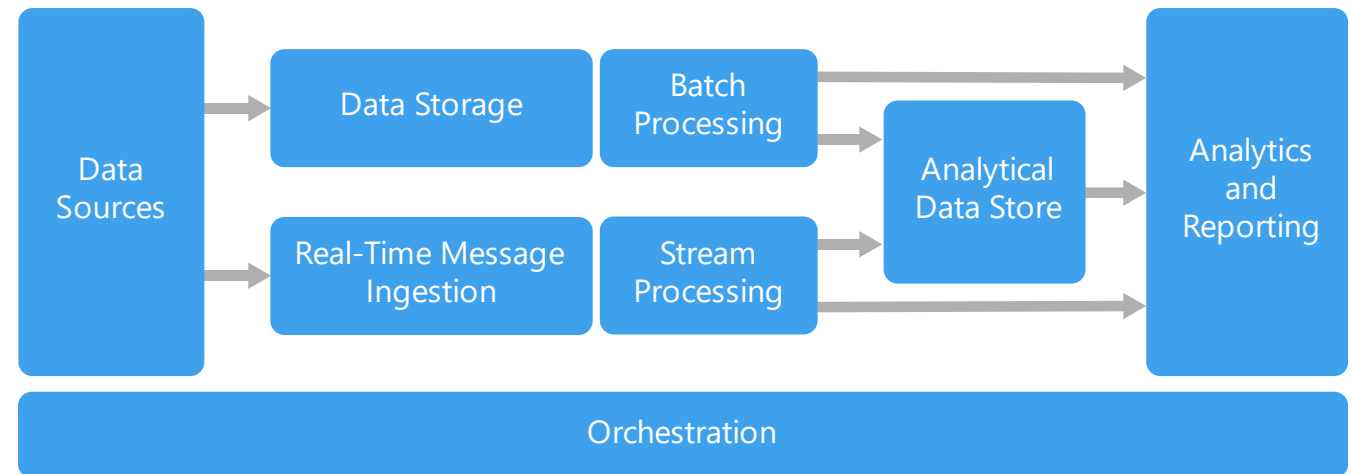
Data Processing on the Cloud

- Data Infrastructures

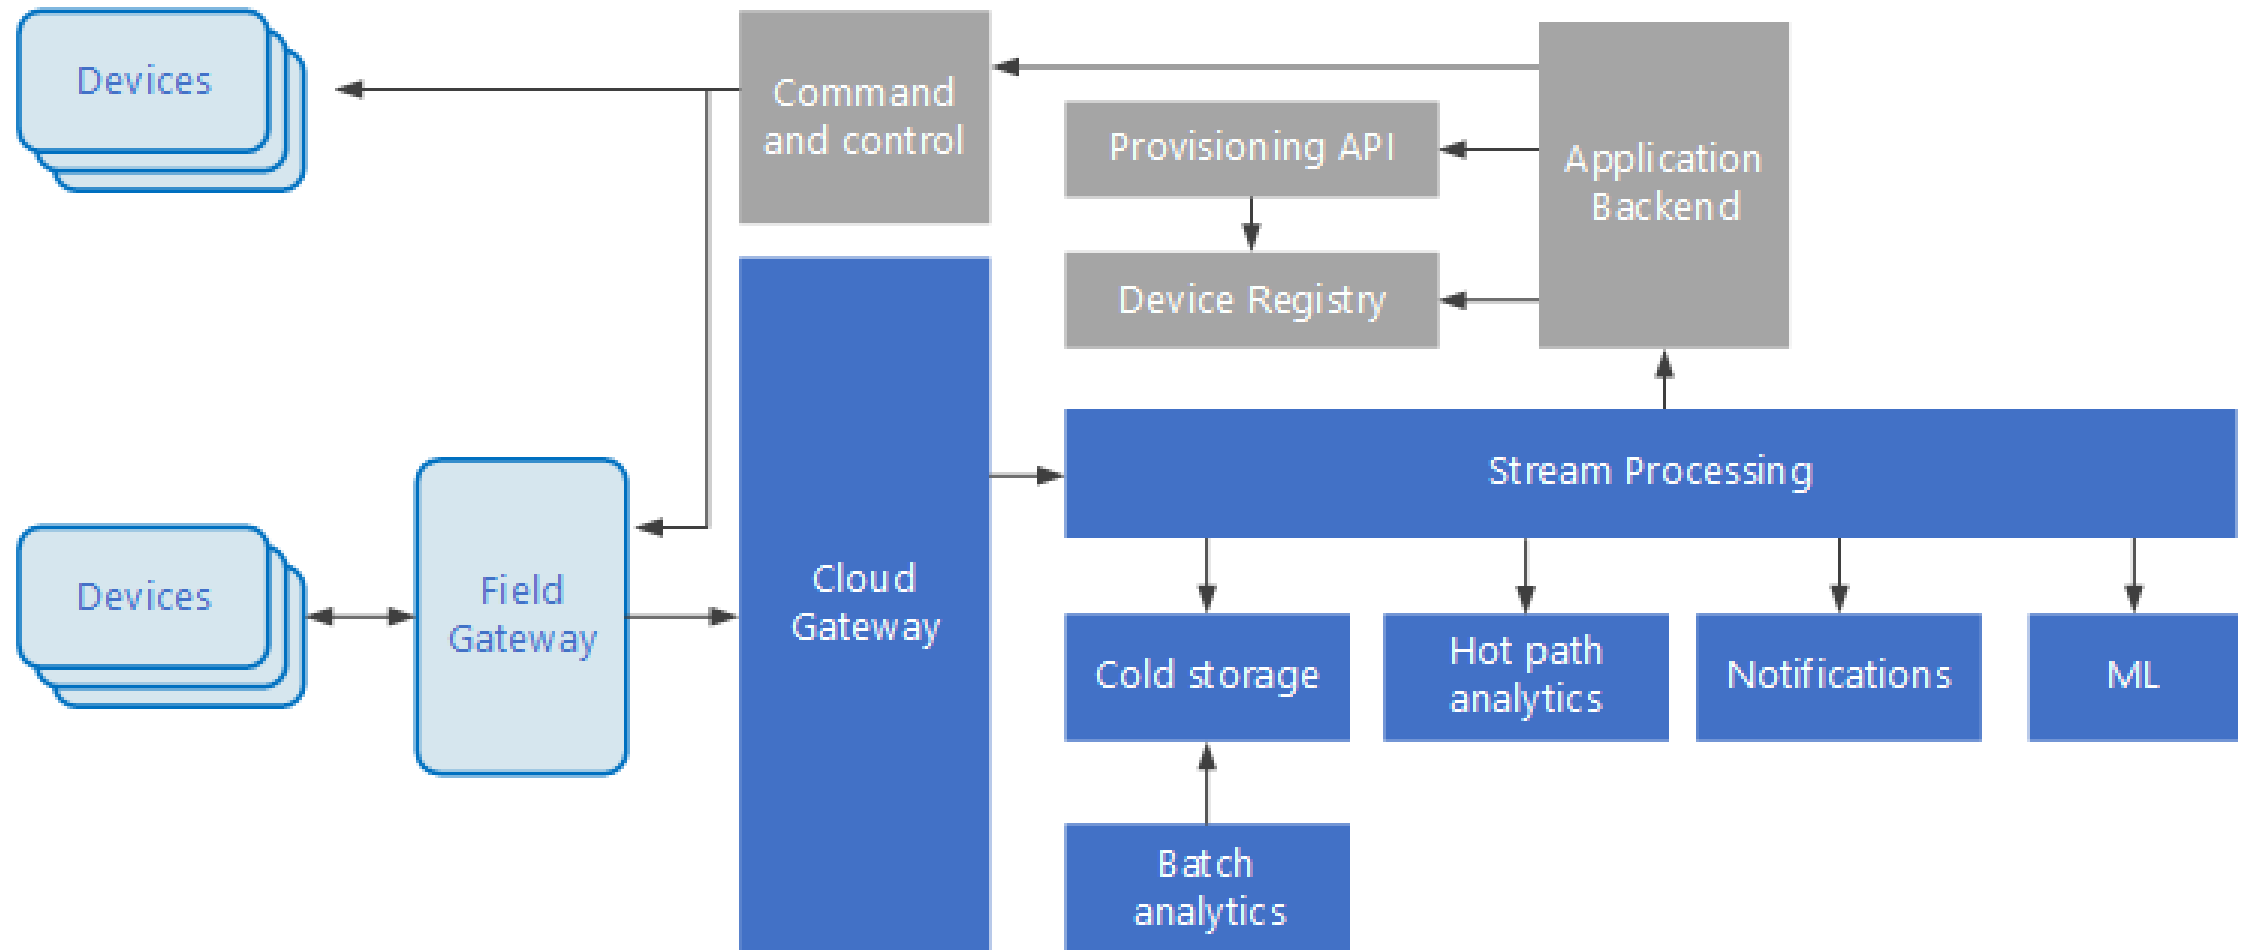
- Batch Processing
- Streaming Data
- Data Integration
- Cloud Native Platforms

- Architecture elements

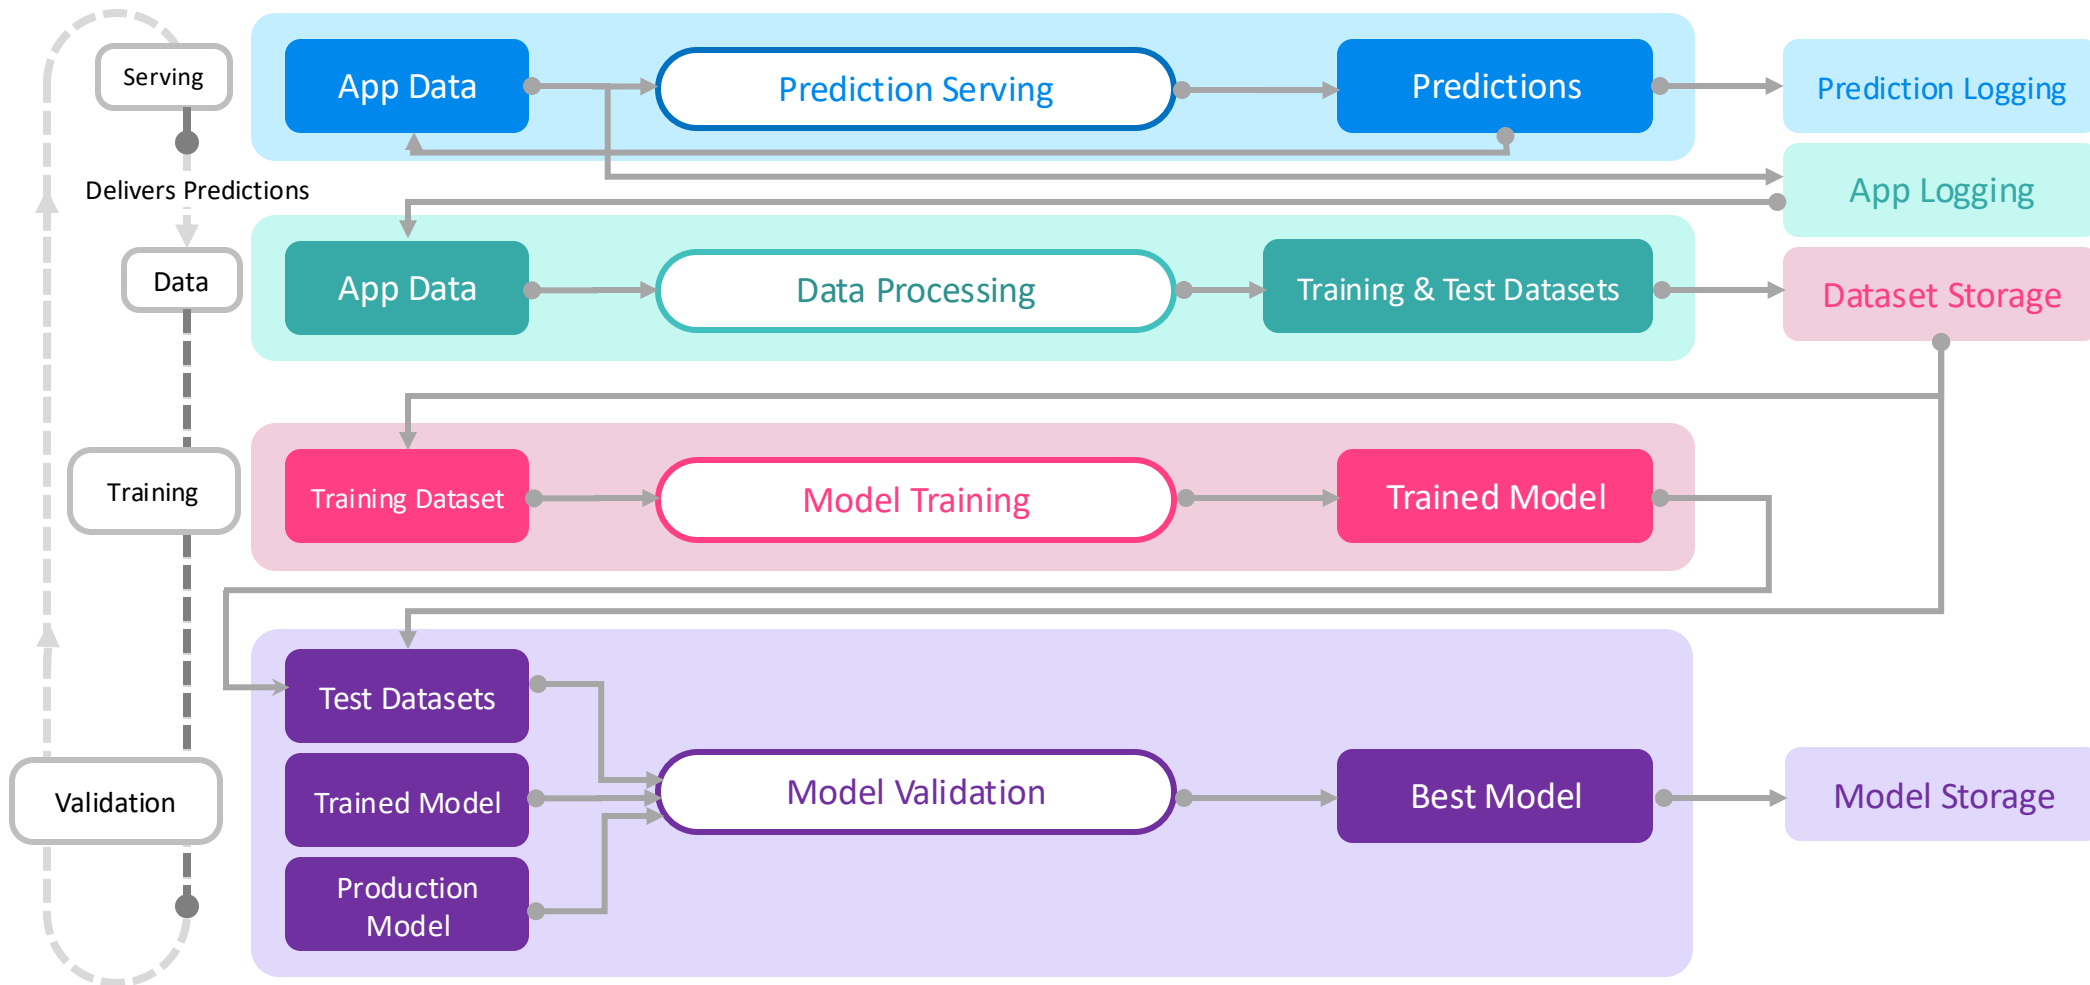
- Ingestion
- Transformation
- Storage



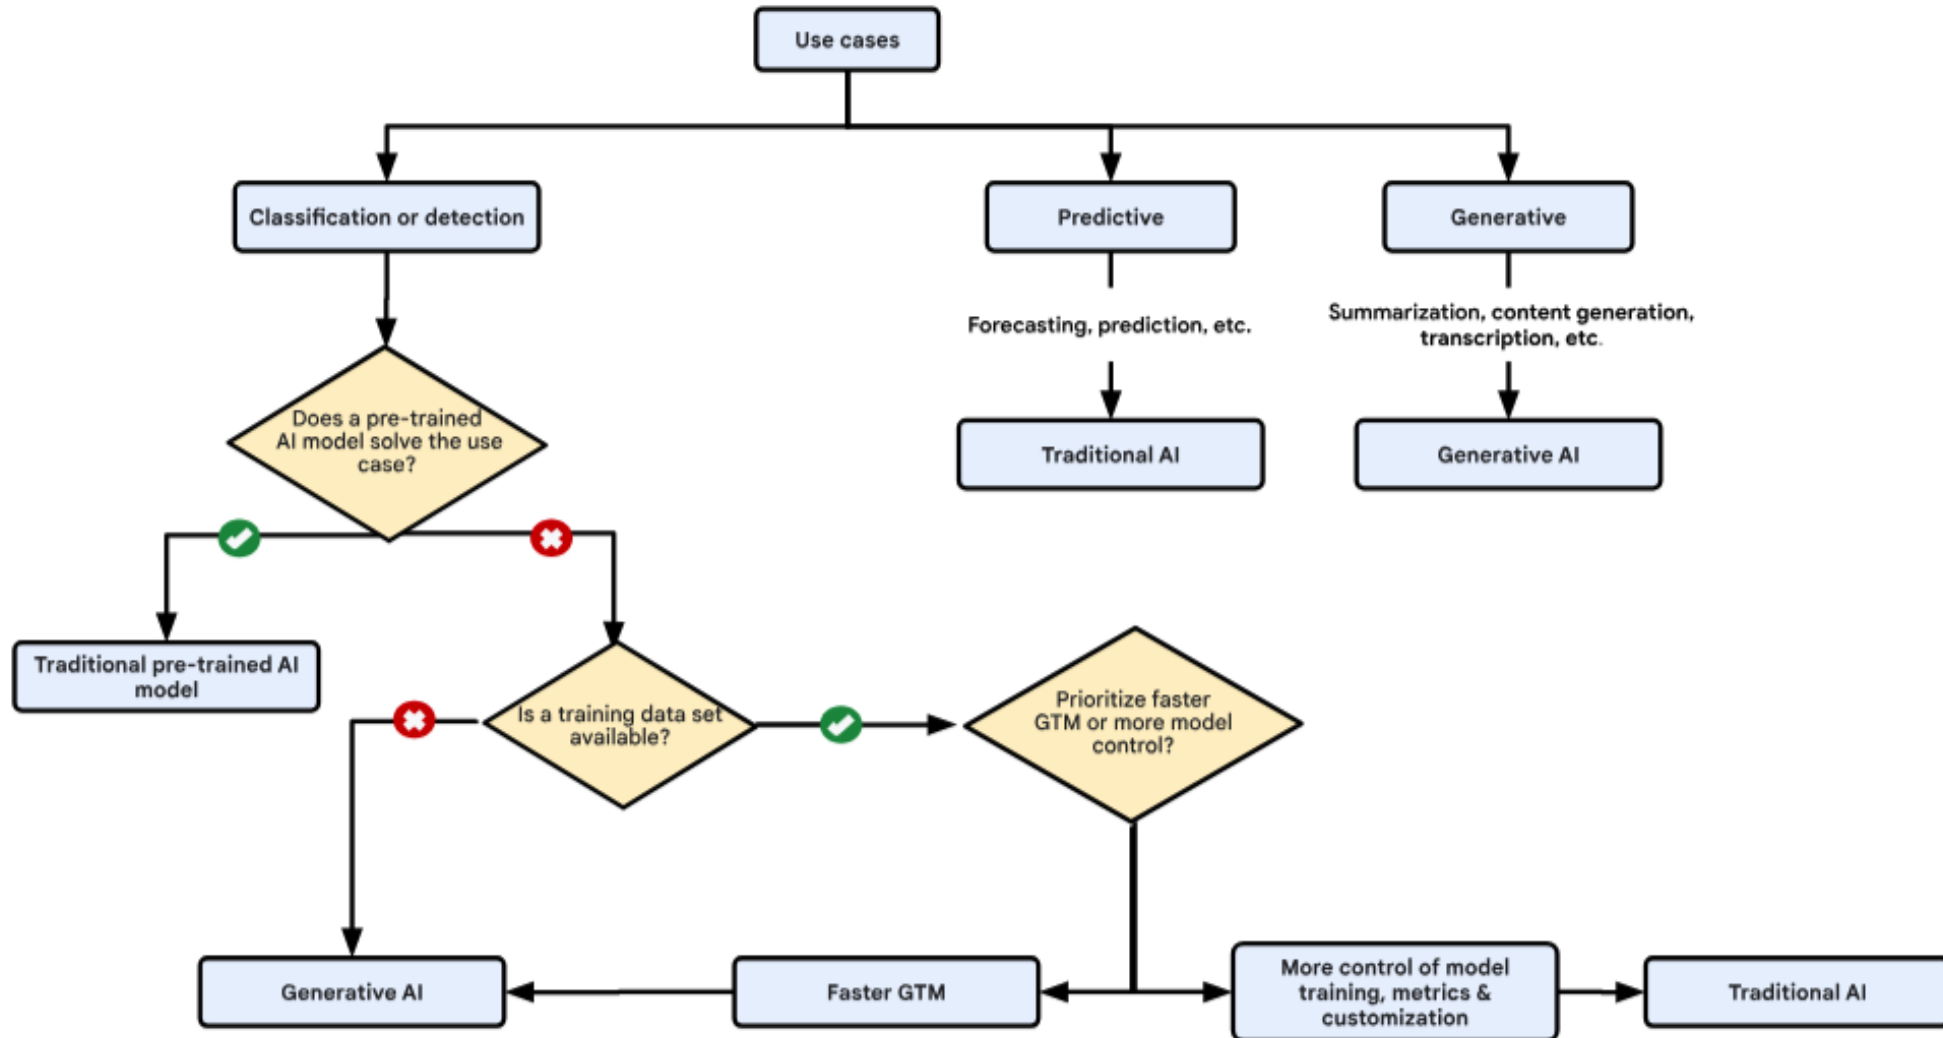
Example: IoT Applications



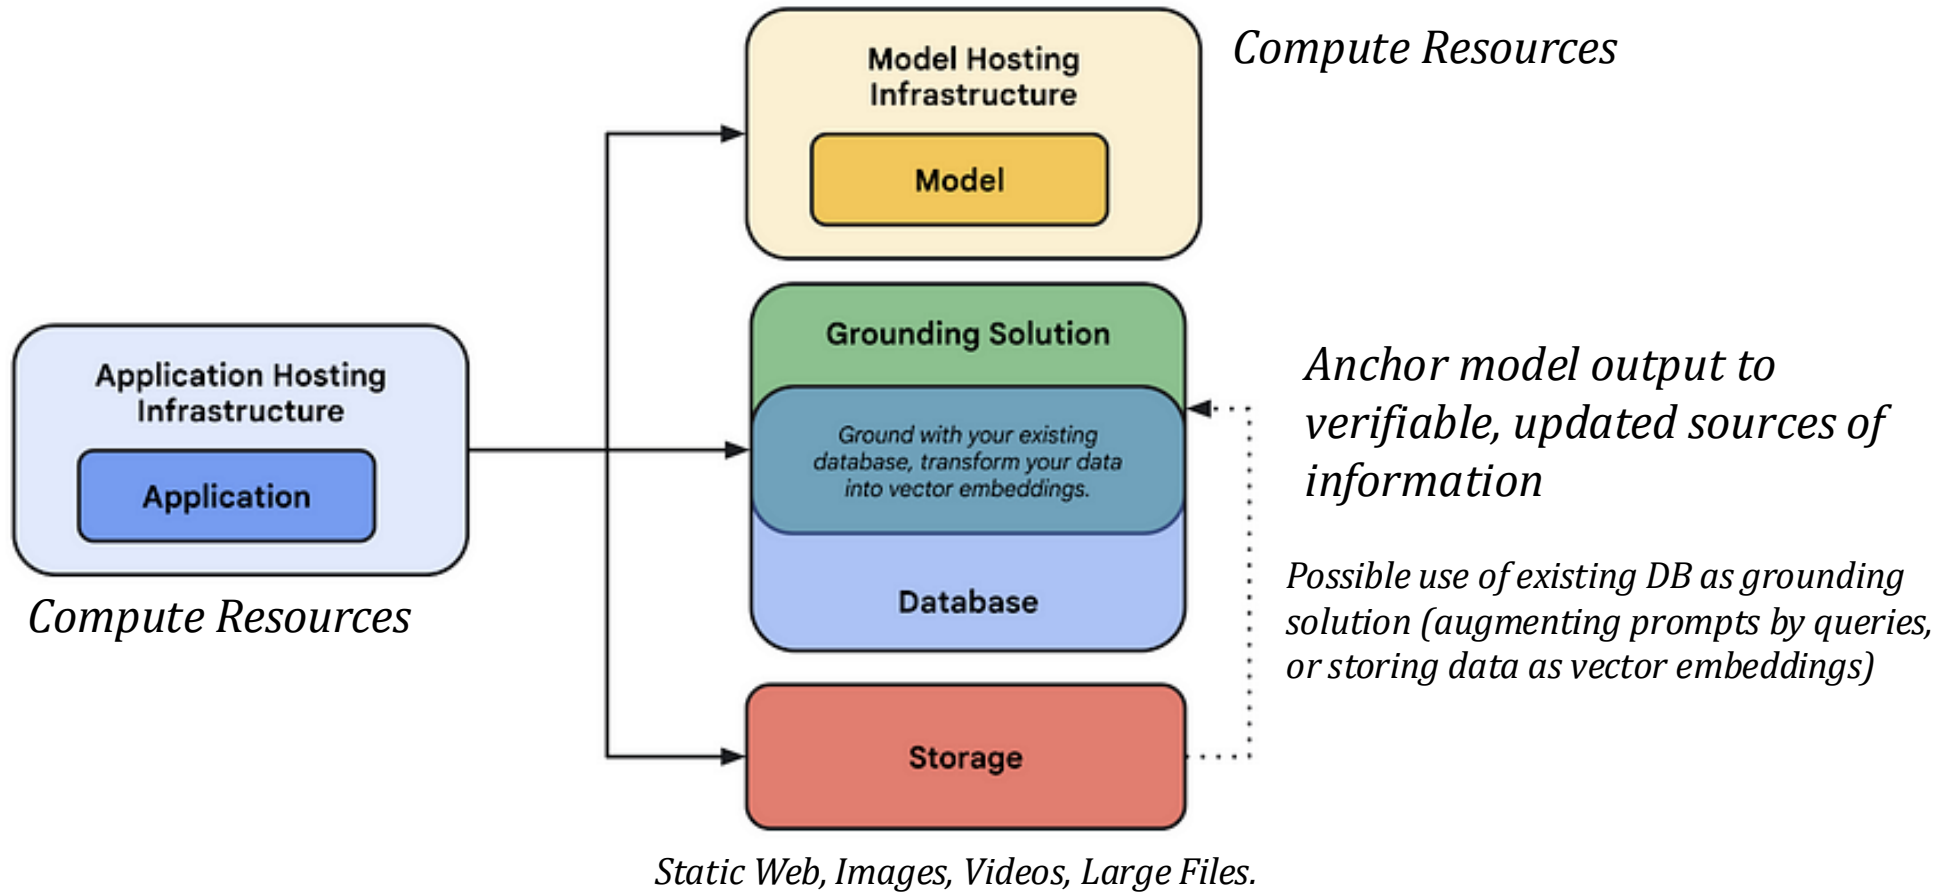
Machine Learning Pipelines



Traditional or Generative



Architecting Applications



Architecting Generative Artificial Intelligence (GenAI) Applications

Summary

- Applications we build using the cloud
 - Popular Use Cases
- What's next?
 - Programming Models & Frameworks
 - Cloud Enablers
 - Resources abstraction
 - Programming abstraction
 - Scalability
 - Parallelization
 - Resource sharing
 - Instance selection
 - Coordination
 - Load-balancing
 - Fault-tolerance

Acknowledgement

The list of resources used in preparation of this slide set are provided on:

<https://canvas.sfu.ca/courses/88212/pages/references>

Pictures and quoted resources are mentioned in each use.

